

Assembly Language Specification

2024 Spring, SWPP

Updates.

1. Architecture Overview

This architecture consists of a single-core CPU and 64-bit memory space.

(1) Registers

- There are 33 64-bit general registers. They are named `r1`, `r2`, ..., `r32`, and `sp`.
- `r1`, `r2`, ..., `r32` are initialized to 0, and `sp` is initialized to 102400.
- Also, there are 16 vector registers. They are named `v1`, `v2`, ..., `v16`.
- `v1`, `v2`, ..., `v16` are initialized to 0.
- Registers are *big-endian*.
- A register can be assigned multiple times (it isn't SSA).

(2) Memory

Loads and stores.

- The memory is accessed via load/store instructions with 64-bit pointers.
- The exact formula for the cost calculation will be described later.

Stack.

- The stack area starts from address 102400, grows downward (-), and is initialized as 0 at the beginning of the program execution.
- You can use `sp` to store the address of the current stack frame, but it is not necessary to do so.

Heap.

- The heap area starts from address 204800 and grows upward (+).
- Heap allocation (`malloc`) initializes the area as zero.
- Accessing an unallocated heap raises an error.
- Accessing the area between [102400, 204800) raises an error.

Global Variables.

- Syntactically, there is no difference between global variables and heap-allocated blocks.
- The project skeleton lowers a global variable to a heap allocation (`malloc` call) at the beginning of the `main()`. So, they are placed at the beginning of the heap area.

(3) Function calls

- Function arguments can be accessed via read-only registers arg1.. arg16.

Calling convention.

- When a call or rcall instruction is executed,
 - r1 ~ r32, v1 ~ v16, sp registers are automatically saved in an invisible space (you don't need to manually spill them).
 - Values of the arguments are automatically assigned to the registers arg1 ~ arg16.
 - Values in registers do not change (they don't get initialized to 0).
- After the call returns, register values are automatically restored to the time before the call.

(4) Cost

- The execution cost of a program can be calculated as 'program-wide instruction execution cost + maximum heap memory usage (in bytes)' * 1024.
- The code size is irrelevant to the total cost.

Memory usage cost.

- The memory usage cost is 1024 times the maximum heap-allocated byte size at any moment.
- For example, the memory usage cost of

```
    r1 = malloc 8
    free r1
    r2 = malloc 8
    free r2
```

is $1024 * 8 = 8192$, because the maximum memory usage is 8 bytes.

Compile time.

- Compile time should be less than 1 minute.

2. Input Program

Structure.

- The source program consists of a single IR file; There is no linking.
- The IR file consists of one or more functions, including the main function.
- A source program only uses i1, i8, i16, i32, i64, array types, pointer types, and vector types.

Function.

- A function can have at most 16 arguments.
- There is no function attribute (e.g. read-only).
- `main()` is never called recursively.

Standard I/O.

- A source program takes input through `read()` calls. `read()` reads an integer and returns it as an i64 value.
- The output of the program is done via `write(i64)` calls. It writes the output as an unsigned integer in a new line.
- `read()` / `write(i64)` calls are connected to the standard input/output.

Misc.

- The test programs will never raise out-of-memory or stack overflow with the given inputs if compiled with the project skeleton.

3. Function & Basic Block

(1) Function

Syntax:

```
start <funcname> <Narg>:  
    ... (basic blocks)  
end <funcname>
```

- A function contains one or more basic blocks.
- <funcname> is a non-empty string consisting of alphabets(a-zA-Z), digits(0-9), underscore(_), hyphen(-), or dot(.).
- <Narg> describes the number of arguments.
- A function's return type is always i64.
- There is no variadic function.
- There is no nested function.

(2) Basic Block

Syntax:

```
<bbname>:  
    ... (instructions)
```

- A basic block consists of one or more instructions.
- A basic block must end with a terminator instruction (see below for more details)
- <bbname> is a non-empty string, starting with a dot(.) and consists of alphabets(a-zA-Z) + digits(0-9) + underscore(_) + hyphen(-) + dot(.).

(3) Comment

Syntax:

```
; <comment>
```

- A comment starts with a semicolon(;).
- Only spaces are allowed before the semicolon in the line.

4. Instructions

Syntax:

```

    op_name <reg1> .. <regN>
    <reg> = op_name <reg1> .. <regN>

```

- <reg>, <regN>, <reg_*> stands for a scalar register.
- <vr>, <vrN>, <vr_*> stands for a vector register.
- <cst>, <cstN>, <cst_*> stands for a nonnegative constant integer ($< 2^{64}$).
- Argument registers (e.g. arg1) cannot be placed at the LHS.

(1) Control Flow

Kind	Syntax	Cost
Return Value - ret is equivalent to ret 0.	ret ret <reg>	1
Unconditional Branch	br <bname>	30
Conditional Branch	br <reg_cond> <>true_bb> <>false_bb>	90 for true_bb 30 for false_bb
Switch Instruction	switch <reg_cond> <cst1> <bb1> .. <default_bb>	60
Function call	call <fname> <reg1> .. <regN> <reg_ret> = call <fname> <reg1> .. <regN>	30
Recursive function call	rcall <reg1> .. <regN> <reg_ret> = rcall <reg1> .. <regN>	10
Assertion An assertion fail is an error. Used for testing	assert_eq <reg1> <reg2> assert_eq <reg> <cst>	0

- ret, br, switch instructions should come at the end of a basic block only.
- <bbname>, <bbN>, <*_bb> is the name of a basic block to jump to.
- <fname> is the name of a function to call.
- rcall calls the function that this instruction resides in.
- br / switch cannot jump to a block in another function.

(2) Memory Allocation/Deallocation

Kind	Syntax	Cost
Heap Allocation	<code><reg_ptr> = malloc <reg_size></code>	150
Deallocation	<code>free <reg_ptr></code>	150

malloc

- The size of `malloc` should be non-zero & a multiple of 8.
- The returned address by `malloc` is a multiple of 8.

free

- `free` deallocates a space associated with the given pointer.
- The pointer passed to `free` should point to the beginning of allocated heap space.

(3) Memory Access

Kind	Syntax	Base Cost
Load	<code><reg> = load <sz> <reg_ptr></code> <code><sz> := 1 2 4 8</code>	Stack area: 30 Heap area: 50
Store	<code>store <sz> <reg> <reg_ptr></code> <code><sz> := 1 2 4 8</code>	Stack area: 30 Heap area: 50
Vector Load	<code><vr> = vload <reg_ptr></code>	Stack area: 60 Heap area: 100
Vector Store	<code>vstore <vr> <reg_ptr></code>	Stack area: 60 Heap area: 100

- Value in `<reg_ptr>` should be multiple of `<sz>`, or multiple of 8 for vector memory address. Unaligned memory access will crash the interpreter.

load

- The `load` instruction reads the data at [`<reg_ptr>` , `<reg_ptr>+<sz>`), zero-extends it to 64 bits, and returns it.
- The memory is *little-endian*. The least significant byte of the value read by `load` is from `<reg_ptr>`, and the most significant byte is from `<reg_ptr>+<sz>-1`.

store

- The `store` instruction truncates the value `<reg>` to an `<sz>*8`-bit integer and writes it at [`<reg_ptr>`, `<reg_ptr>+<sz>`).

vector load

- The `vload` instruction reads the data at [`<reg_ptr>` , `<reg_ptr>+32`) and writes it into destination `<vr>`

vector store

- The `vstore` instruction writes the value of `<vr>` at [`<reg_ptr>`, `<reg_ptr>+32`).

(4) Scalar Arithmetic

Kind	Name	Cost
Integer Multiplication/Division	$\langle \text{reg} \rangle = \text{udiv } \langle \text{reg1} \rangle \langle \text{reg2} \rangle \langle \text{bw} \rangle$ $\langle \text{reg} \rangle = \text{sdiv } \langle \text{reg1} \rangle \langle \text{reg2} \rangle \langle \text{bw} \rangle$ $\langle \text{reg} \rangle = \text{urem } \langle \text{reg1} \rangle \langle \text{reg2} \rangle \langle \text{bw} \rangle$ $\langle \text{reg} \rangle = \text{srem } \langle \text{reg1} \rangle \langle \text{reg2} \rangle \langle \text{bw} \rangle$ $\langle \text{reg} \rangle = \text{mul } \langle \text{reg1} \rangle \langle \text{reg2} \rangle \langle \text{bw} \rangle$ $\langle \text{bw} \rangle := 1 8 16 32 64$	1
Integer Shift - shl: shift-left - lshr: logical shift-right - ashr: arithmetic shift-right	$\langle \text{reg} \rangle = \text{shl } \langle \text{reg1} \rangle \langle \text{reg2} \rangle \langle \text{bw} \rangle$ $\langle \text{reg} \rangle = \text{lshr } \langle \text{reg1} \rangle \langle \text{reg2} \rangle \langle \text{bw} \rangle$ $\langle \text{reg} \rangle = \text{ashr } \langle \text{reg1} \rangle \langle \text{reg2} \rangle \langle \text{bw} \rangle$ $\langle \text{bw} \rangle := 1 8 16 32 64$	4
Bitwise Operation	$\langle \text{reg} \rangle = \text{and } \langle \text{reg1} \rangle \langle \text{reg2} \rangle \langle \text{bw} \rangle$ $\langle \text{reg} \rangle = \text{or } \langle \text{reg1} \rangle \langle \text{reg2} \rangle \langle \text{bw} \rangle$ $\langle \text{reg} \rangle = \text{xor } \langle \text{reg1} \rangle \langle \text{reg2} \rangle \langle \text{bw} \rangle$ $\langle \text{bw} \rangle := 1 8 16 32 64$	4
Integer Add/Sub	$\langle \text{reg} \rangle = \text{add } \langle \text{reg1} \rangle \langle \text{reg2} \rangle \langle \text{bw} \rangle$ $\langle \text{reg} \rangle = \text{sub } \langle \text{reg1} \rangle \langle \text{reg2} \rangle \langle \text{bw} \rangle$ $\langle \text{bw} \rangle := 1 8 16 32 64$	5
Integer increment $\langle \text{reg} \rangle = \langle \text{reg} \rangle + 1$	$\langle \text{reg} \rangle = \text{incr } \langle \text{reg1} \rangle \langle \text{bw} \rangle$ $\langle \text{bw} \rangle := 1 8 16 32 64$	1
Integer decrement $\langle \text{reg} \rangle = \langle \text{reg} \rangle - 1$	$\langle \text{reg} \rangle = \text{decr } \langle \text{reg1} \rangle \langle \text{bw} \rangle$ $\langle \text{bw} \rangle := 1 8 16 32 64$	1
Comparison - $\langle \text{cond} \rangle$ is equivalent to the cond of LLVM IR's icmp	$\langle \text{reg} \rangle = \text{icmp } \langle \text{cond} \rangle \langle \text{reg1} \rangle \langle \text{reg2} \rangle \langle \text{bw} \rangle$ $\langle \text{bw} \rangle := 1 8 16 32 64$	1
Ternary operation	$\langle \text{reg} \rangle = \text{select } \langle \text{reg_cond} \rangle$ $\quad \quad \quad \langle \text{reg_true} \rangle \langle \text{reg_false} \rangle$	1
Move constant	$\langle \text{reg} \rangle = \text{const } \langle \text{cst} \rangle$	1

- For integer arithmetic and comparison operations, $\langle \text{bw} \rangle$ is the size of bitwidth of inputs that the operation assumes. For example, ``ashr 511 2 8`` takes the lowest 8-bits from inputs (which is $255 = -1$), performs arithmetic right shift, and zero-extends it to 64 bits. So, its result is 255.
- For select, if $\langle \text{reg_cond} \rangle$ is not one of 0 or 1, the interpreter will crash.

(5) Elementwise Vector Arithmetic

Kind	Name	Cost
Integer Multiplication/Division	$\langle vr \rangle = vdiv \ \langle vr1 \rangle \ \langle vr2 \rangle \ \langle bw \rangle$ $\langle vr \rangle = vsdiv \ \langle vr1 \rangle \ \langle vr2 \rangle \ \langle bw \rangle$ $\langle vr \rangle = vurem \ \langle vr1 \rangle \ \langle vr2 \rangle \ \langle bw \rangle$ $\langle vr \rangle = vsrem \ \langle vr1 \rangle \ \langle vr2 \rangle \ \langle bw \rangle$ $\langle vr \rangle = vmul \ \langle vr1 \rangle \ \langle vr2 \rangle \ \langle bw \rangle$ $\langle bw \rangle := 32 64$	2
Integer Shift - shl: shift-left - lshr: logical shift-right - ashr: arithmetic shift-right	$\langle vr \rangle = vshl \ \langle vr1 \rangle \ \langle vr2 \rangle \ \langle bw \rangle$ $\langle vr \rangle = vlshr \ \langle vr1 \rangle \ \langle vr2 \rangle \ \langle bw \rangle$ $\langle vr \rangle = vashr \ \langle vr1 \rangle \ \langle vr2 \rangle \ \langle bw \rangle$ $\langle bw \rangle := 32 64$	8
Bitwise Operation	$\langle vr \rangle = vand \ \langle vr1 \rangle \ \langle vr2 \rangle \ \langle bw \rangle$ $\langle vr \rangle = vor \ \langle vr1 \rangle \ \langle vr2 \rangle \ \langle bw \rangle$ $\langle vr \rangle = vxor \ \langle vr1 \rangle \ \langle vr2 \rangle \ \langle bw \rangle$ $\langle bw \rangle := 32 64$	8
Integer Add/Sub	$\langle vr \rangle = vadd \ \langle vr1 \rangle \ \langle vr2 \rangle \ \langle bw \rangle$ $\langle vr \rangle = vsub \ \langle vr1 \rangle \ \langle vr2 \rangle \ \langle bw \rangle$ $\langle bw \rangle := 32 64$	10
Integer Increment	$\langle vr \rangle = vincr \ \langle vr1 \rangle \ \langle bw \rangle$ $\langle bw \rangle := 32 64$	2
integer Decrement	$\langle vr \rangle = vdecr \ \langle vr1 \rangle \ \langle bw \rangle$ $\langle bw \rangle := 32 64$	2
Comparison - $\langle cond \rangle$ is equivalent to the cond of LLVM IR's icmp	$\langle vr \rangle = vicmp \ \langle cond \rangle \ \langle vr1 \rangle \ \langle vr2 \rangle \ \langle bw \rangle$ $\langle bw \rangle := 32 64$	2
Ternary operation	$\langle vr \rangle = vselect \ \langle vr_cond \rangle$ $\quad \quad \quad \langle vr_true \rangle \ \langle vr_false \rangle \ \langle bw \rangle$ $\langle bw \rangle := 32 64$	2

- All arithmetic operations described above are executed element-wise.

They execute scalar arithmetic operations between elements of the same index from two vector registers, or repeat the operation for every element if the operation is unary.

- For ternary operation, if each element in $\langle vr_cond \rangle$ is not one of 0 or 1, the interpreter will crash.

(6) Parallel Vector Arithmetic

Kind	Name	Cost
Integer Multiplication/Division	<code><vr> = vpudiv <vr1> <vr2> <bw></code> <code><vr> = vpsdiv <vr1> <vr2> <bw></code> <code><vr> = vpurem <vr1> <vr2> <bw></code> <code><vr> = vpsrem <vr1> <vr2> <bw></code> <code><vr> = vpmul <vr1> <vr2> <bw></code> <code><bw> := 32 64</code>	2
Bitwise Operation	<code><vr> = vpand <vr1> <vr2> <bw></code> <code><vr> = vpor <vr1> <vr2> <bw></code> <code><vr> = vpxor <vr1> <vr2> <bw></code> <code><bw> := 32 64</code>	8
Integer Add/Sub	<code><vr> = vpadd <vr1> <vr2> <bw></code> <code><vr> = vpsub <vr1> <vr2> <bw></code> <code><bw> := 32 64</code>	10
Comparison - <cond> is equivalent to the cond of LLVM IR's icmp	<code><vr> = vpicmp <cond> <vr1> <vr2> <bw></code> <code><bw> := 32 64</code>	2
Ternary operation	<code><vr> = vpselect <vr_cond> <vr1> <vr2> <bw></code> <code><bw> := 32 64</code>	2

- All arithmetic operations described above are executed in a 'parallel' manner. They execute scalar arithmetic operations between each even-index element and adjacent odd-index element on the right.

For example, `v3 = vpadd v1 v2 64` executes the following:

`v3[0] <- v1[0] + v1[1]`

`v3[1] <- v1[2] + v1[3]`

`v3[2] <- v2[0] + v2[1]`

`v3[3] <- v2[2] + v2[3]`

- For ternary operation, if each element in <vr_cond> is not one of 0 or 1, the interpreter will crash.

(7) Vector Manipulation

Kind	Name	Cost
Broadcast	$\langle vr \rangle = \text{vbcst } \langle reg \rangle \langle bw \rangle$ $\langle bw \rangle := 32 64$	4
Extract	$\langle reg \rangle = \text{vextct } \langle vr1 \rangle \langle reg_idx \rangle \langle bw \rangle$ $\langle bw \rangle := 32 64$	4
Update	$\langle vr \rangle = \text{vupdate } \langle vr1 \rangle \langle reg \rangle \langle reg_idx \rangle \langle bw \rangle$ $\langle bw \rangle := 32 64$	4

- For extract and update, $\langle reg_idx \rangle$ must be less than 8 if $\langle bw \rangle$ is 32, and less than 4 if $\langle bw \rangle$ is 64. Otherwise the interpreter will crash.
- Broadcast writes $\langle reg \rangle$ to every element in $\langle vr \rangle$, with each element size assumed as $\langle bw \rangle$.
- Extract reads the $\langle reg_idx \rangle$ th element of $\langle vr1 \rangle$ and write the result at $\langle reg \rangle$
- Update changes the $\langle reg_idx \rangle$ th element of $\langle vr1 \rangle$ to $\langle reg \rangle$ and write the result at $\langle vr \rangle$